Refactor Udagram App into Microservices and Deploy

Tags

Screenshot of DockerHub shows the images

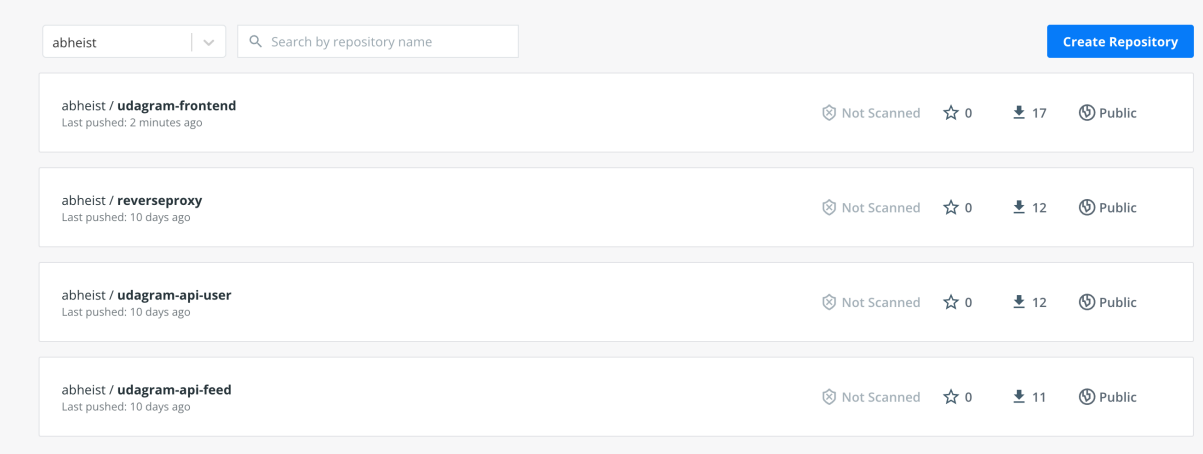

Screenshot of the Travis CI interface shows a successful build and deploy job

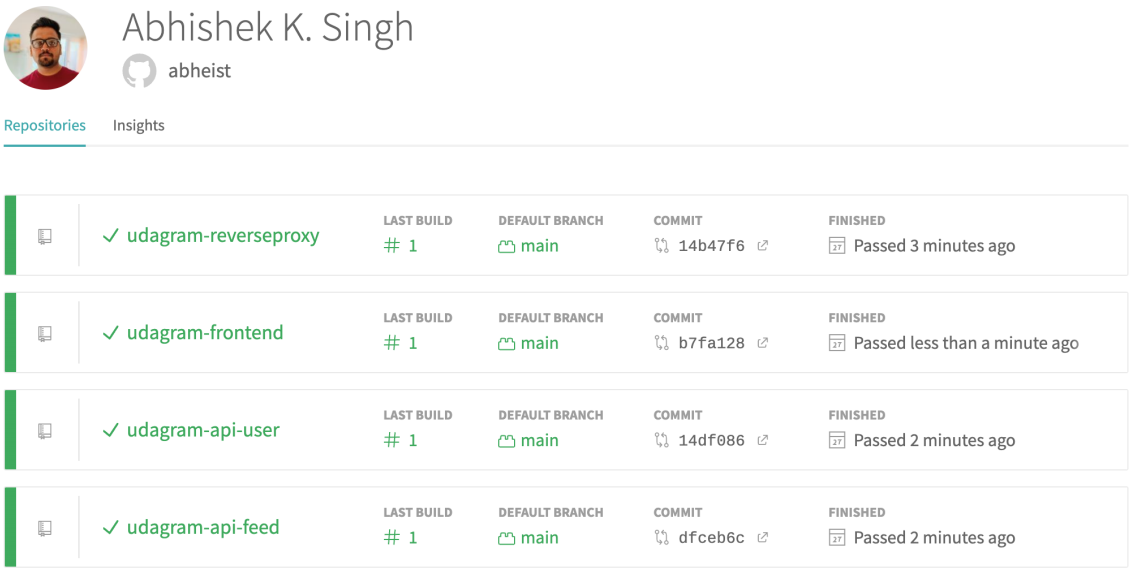

A screenshots of `kubectl` commands show the Frontend and API projects deployed in Kubernetes.

```
cd0354-monolith-to-microservices-project on  main [+]
➜ kubectl get deployments.apps
NAME           READY   UP-TO-DATE   AVAILABLE   AGE
backend-feed   3/3     3            3           3h35m
backend-user   3/3     3            3           3h34m
frontend       1/1     1            1           93m
reverseproxy   1/1     1            1           3h34m
```

```
cd0354-monolith-to-microservices-project on  main [+]
➜ kubectl get pods
NAME                            READY   STATUS    RESTARTS   AGE
backend-feed-66f78986bf-f6hzq   1/1     Running   0          59m
backend-feed-66f78986bf-pjt6s   1/1     Running   0          59m
backend-feed-66f78986bf-xwh5q   1/1     Running   0          3h35m
backend-user-8699fdfff4-g5pt7   1/1     Running   0          59m
backend-user-8699fdfff4-jqnsh   1/1     Running   0          3h35m
backend-user-8699fdfff4-klmhz   1/1     Running   0          59m
frontend-6cf49d8677-nm7zr       1/1     Running   0          68m
reverseproxy-6498f8f6db-6brtl   1/1     Running   0          3h34m
```

```
cd0354-monolith-to-microservices-project on  main [+]
➜ kubectl describe services
Name:              backend-feed
Namespace:         default
Labels:            service=backend-feed
Annotations:       <none>
Selector:          service=backend-feed
Type:              ClusterIP
IP Family Policy:  SingleStack
IP Families:       IPv4
IP:                10.100.35.146
IPs:               10.100.35.146
Port:              8080  8080/TCP
TargetPort:        8080/TCP
Endpoints:         172.31.23.247:8080,172.31.85.100:8080,172.31.92.186:8080
Session Affinity:  None
Events:            <none>


Name:              backend-user
```

```
Namespace:          default
Labels:             service=backend-user
Annotations:        <none>
Selector:           service=backend-user
Type:               ClusterIP
IP Family Policy:   SingleStack
IP Families:        IPv4
IP:                 10.100.127.119
IPs:                10.100.127.119
Port:               8080  8080/TCP
TargetPort:         8080/TCP
Endpoints:          172.31.18.219:8080,172.31.84.82:8080,172.31.93.85:8080
Session Affinity:   None
Events:             <none>


Name:               frontend
Namespace:          default
Labels:             service=frontend
Annotations:        <none>
Selector:           service=frontend
Type:               ClusterIP
IP Family Policy:   SingleStack
IP Families:        IPv4
IP:                 10.100.50.0
IPs:                10.100.50.0
Port:               8100  8100/TCP
TargetPort:         80/TCP
Endpoints:          172.31.27.79:80
Session Affinity:   None
Events:             <none>


Name:               kubernetes
Namespace:          default
Labels:             component=apiserver
                    provider=kubernetes
Annotations:        <none>
Selector:           <none>
Type:               ClusterIP
IP Family Policy:   SingleStack
IP Families:        IPv4
IP:                 10.100.0.1
IPs:                10.100.0.1
Port:               https  443/TCP
TargetPort:         443/TCP
Endpoints:          172.31.3.231:443,172.31.86.139:443
Session Affinity:   None
Events:             <none>


Name:                     publicfrontend
Namespace:                default
Labels:                   service=frontend
Annotations:              <none>
Selector:                 service=frontend
Type:                     LoadBalancer
IP Family Policy:         SingleStack
```

```
IP Families:              IPv4
IP:                       10.100.73.65
IPs:                      10.100.73.65
LoadBalancer Ingress:     af6cd9566ed7e44489ffeef8c2f2c391-762354092.us-east-1.elb.ama
zonaws.com
Port:                     <unset>  80/TCP
TargetPort:               80/TCP
NodePort:                 <unset>  30859/TCP
Endpoints:                172.31.27.79:80
Session Affinity:         None
External Traffic Policy:  Cluster
Events:                   <none>


Name:                     publicreverseproxy
Namespace:                default
Labels:                   service=reverseproxy
Annotations:              <none>
Selector:                 service=reverseproxy
Type:                     LoadBalancer
IP Family Policy:         SingleStack
IP Families:              IPv4
IP:                       10.100.97.220
IPs:                      10.100.97.220
LoadBalancer Ingress:     a2eadaae92d784276ad3cd7dddae569d-1975764867.us-east-1.elb.am
azonaws.com
Port:                     <unset>  8080/TCP
TargetPort:               8080/TCP
NodePort:                 <unset>  31720/TCP
Endpoints:                172.31.87.109:8080
Session Affinity:         None
External Traffic Policy:  Cluster
Events:                   <none>


Name:              reverseproxy
Namespace:         default
Labels:            service=reverseproxy
Annotations:       <none>
Selector:          service=reverseproxy
Type:              ClusterIP
IP Family Policy:  SingleStack
IP Families:       IPv4
IP:                10.100.220.100
IPs:               10.100.220.100
Port:              8080  8080/TCP
TargetPort:        8080/TCP
Endpoints:         172.31.87.109:8080
Session Affinity:  None
Events:            <none>
```

HPA - replicas

```
cd0354-monolith-to-microservices-project on  main [+]
➜ kubectl get horizontalpodautoscalers.autoscaling
NAME           REFERENCE                  TARGETS         MINPODS   MAXPODS   REPLICAS   AGE
backend-feed   Deployment/backend-feed   <unknown>/70%   3         5         3          61m
backend-user   Deployment/backend-user   <unknown>/70%   3         5         3          61m
```

```
Name:                                                  backend-feed
Namespace:                                             default
Labels:                                                <none>
Annotations:                                           <none>
CreationTimestamp:                                     Mon, 04 Jul 2022 22:14:52 -0300
Reference:                                             Deployment/backend-feed
Metrics:                                               ( current / target )
  resource cpu on pods  (as a percentage of request):  <unknown> / 70%
Min replicas:                                          3
Max replicas:                                          5
Deployment pods:                                       3 current / 3 desired
Conditions:
  Type           Status  Reason                   Message
  ----           ------  ------                   -------
  AbleToScale    True    SucceededGetScale        the HPA controller was able to get t
he target's current scale
  ScalingActive  False   FailedGetResourceMetric  the HPA was unable to compute the re
plica count: failed to get cpu utilization: unable to get metrics for resource cpu: un
able to fetch metrics from resource metrics API: the server could not find the request
ed resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                        Age                 From
Message
  ----     ------                        ----                ----
-------
  Warning  FailedComputeMetricsReplicas  58m (x12 over 60m)  horizontal-pod-autoscale
r  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilizatio
n: unable to get metrics for resource cpu: unable to fetch metrics from resource metri
cs API: the server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedGetResourceMetric       60s (x240 over 60m)  horizontal-pod-autoscale
r  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fe
tch metrics from resource metrics API: the server could not find the requested resourc
e (get pods.metrics.k8s.io)


Name:                                                  backend-user
Namespace:                                             default
Labels:                                                <none>
Annotations:                                           <none>
CreationTimestamp:                                     Mon, 04 Jul 2022 22:14:41 -0300
Reference:                                             Deployment/backend-user
Metrics:                                               ( current / target )
  resource cpu on pods  (as a percentage of request):  <unknown> / 70%
Min replicas:                                          3
Max replicas:                                          5
Deployment pods:                                       3 current / 3 desired
Conditions:
  Type           Status  Reason                   Message
  ----           ------  ------                   -------
  AbleToScale    True    SucceededGetScale        the HPA controller was able to get t
```

```
he target's current scale
  ScalingActive  False   FailedGetResourceMetric  the HPA was unable to compute the re
plica count: failed to get cpu utilization: unable to get metrics for resource cpu: un
able to fetch metrics from resource metrics API: the server could not find the request
ed resource (get pods.metrics.k8s.io)
Events:
  Type     Reason                        Age                  From
Message
  ----     ------                        ----                 ----
-------
  Warning  FailedComputeMetricsReplicas  58m (x12 over 61m)   horizontal-pod-autoscale
r  invalid metrics (1 invalid out of 1), first error is: failed to get cpu utilizatio
n: unable to get metrics for resource cpu: unable to fetch metrics from resource metri
cs API: the server could not find the requested resource (get pods.metrics.k8s.io)
  Warning  FailedGetResourceMetric       75s (x240 over 61m)  horizontal-pod-autoscale
r  failed to get cpu utilization: unable to get metrics for resource cpu: unable to fe
tch metrics from resource metrics API: the server could not find the requested resourc
e (get pods.metrics.k8s.io)
```

Screenshot of one of the backend API pod logs